

Keras -- MLPs on MNIST

```
In [0]: # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal #or xaiver/Hae normilization

# Importing libraries
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np
import time
```

```
In [0]: '''%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()'''

# this function is used draw Categorical Crossentropy Loss VS No. of epochs plot
def plt_dynamic(x, vy, ty):
    plt.figure(figsize=(10,5))
    plt.plot(x, vy, 'b', label="Validation Loss")
    plt.plot(x, ty, 'r', label="Train Loss")
    plt.xlabel('Epochs')
    plt.ylabel('Categorical Crossentropy Loss')
    plt.title('\n Categorical Crossentropy Loss VS Epochs')
    plt.legend()
    plt.grid()
    plt.show()
```

```
In [0]: # the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
In [0]: print("Number of training examples :", X_train.shape[0], "and each image is of  
shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))  
print("Number of training examples :", X_test.shape[0], "and each image is of  
shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)

Number of training examples : 10000 and each image is of shape (28, 28)

```
In [0]: # if you observe the input shape its 2 dimensional vector  
# for each image we have a (28*28) vector  
# we will convert the (28*28) vector into single dimensional vector of 1 * 784  
  
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])  
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
In [0]: # after converting the input images from 3d to 2d vectors  
  
print("Number of training examples :", X_train.shape[0], "and each image is of  
shape (%d)"%(X_train.shape[1]))  
print("Number of training examples :", X_test.shape[0], "and each image is of  
shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)

Number of training examples : 10000 and each image is of shape (784)

```
In [0]: # An example data point
#pixel value near to 255 is black and near to 0 is white. in middel is gray
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175 26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30 36 94 154
170 253 253 253 253 253 225 172 253 242 195 64  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251 93 82
 82 56 39  0  0  0  0  0  0  0  0  0  0  0  0 18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205 11  0 43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253 90  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253 70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0 35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119 25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0 45 186 253 253 150 27  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 16 93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0 249 253 249 64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0 46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0 39 148 229 253 253 253 250 182  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0 24 114 221 253 253 253
253 201 78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0 23 66 213 253 253 253 253 198 81  2  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0 18 171 219 253 253 253 195
80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
55 172 226 253 253 253 253 244 133 11  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132 16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0]
```

```
In [0]: # if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train = X_train/255
X_test = X_test/255
```

```
In [0]: # example data point after normlizing  
print(X_train[0])
```

6/61

0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.15294118	0.58039216
0.89803922	0.99215686	0.99215686	0.99215686	0.98039216	0.71372549
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.09411765	0.44705882	0.86666667	0.99215686	0.99215686	0.99215686
0.99215686	0.78823529	0.30588235	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.09019608	0.25882353	0.83529412	0.99215686
0.99215686	0.99215686	0.99215686	0.77647059	0.31764706	0.00784314
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.07058824	0.67058824
0.85882353	0.99215686	0.99215686	0.99215686	0.99215686	0.76470588
0.31372549	0.03529412	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.21568627	0.6745098	0.88627451	0.99215686	0.99215686	0.99215686
0.99215686	0.95686275	0.52156863	0.04313725	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.53333333	0.99215686
0.99215686	0.99215686	0.83137255	0.52941176	0.51764706	0.0627451

[illegible]

```
In [0]: # here we are having a class number for each image
print("Class label of first image :", y_train[0])

# Lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0,
0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10) #one hot incoding
Y_test = np_utils.to_categorical(y_test, 10)#one hot incoding

print("After converting the output into a vector : ",Y_train[0])
```

Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

Simple Softmax classifier with optimizer='sgd'


```

In [0]: # https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(W.T. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument supported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

```

```
from keras.models import Sequential
from keras.layers import Dense, Activation
```

```
In [0]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

```
In [0]: # start building a model
model = Sequential()#output of one model goes to input of second model

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

# output_dim represents the number of nodes needed in that layer
# here we have 10 nodes

model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))
```

```
In [0]: # Before training a model, you need to configure the learning process, which is done via the compile method

# It receives three arguments:
# An optimizer. This could be the string identifier of an existing optimizer ,
# https://keras.io/optimizers/
# A loss function. This is the objective that the model will try to minimize.,
# https://keras.io/losses/
# A list of metrics. For any classification problem you will want to set this
# to metrics=['accuracy']. https://keras.io/metrics/

# Note: when using the categorical_crossentropy loss, your targets should be in
# categorical format
# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional
# vector that is all-zeros except
# for a 1 at the index corresponding to the class of the sample).

# that is why we converted out labels into vectors

model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None,
# validation_split=0.0,
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None,
# initial_epoch=0, steps_per_epoch=None,
# validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on
# a dataset).

# it returns A History object. Its History.history attribute is a record of training
# loss values and
# metrics values at successive epochs, as well as validation loss values and validation
# metrics values (if applicable).

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
                    verbose=1, validation_data=(X_test, Y_test))
```

W0623 18:21:52.679665 140469971335040 deprecation_wrapper.py:119] From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:790: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

W0623 18:21:52.716533 140469971335040 deprecation_wrapper.py:119] From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3295: The name tf.log is deprecated. Please use tf.math.log instead.

W0623 18:21:52.828167 140469971335040 deprecation.py:323] From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_grad.py:1250: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

W0623 18:21:52.872838 140469971335040 deprecation_wrapper.py:119] From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:986: The name tf.assign_add is deprecated. Please use tf.compat.v1.assign_add instead.

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 82us/step - loss: 1.2807 -
acc: 0.7030 - val_loss: 0.8106 - val_acc: 0.8351

Epoch 2/20

60000/60000 [=====] - 1s 24us/step - loss: 0.7120 -
acc: 0.8460 - val_loss: 0.6045 - val_acc: 0.8635

Epoch 3/20

60000/60000 [=====] - 1s 24us/step - loss: 0.5833 -
acc: 0.8628 - val_loss: 0.5229 - val_acc: 0.8749

Epoch 4/20

60000/60000 [=====] - 1s 24us/step - loss: 0.5223 -
acc: 0.8710 - val_loss: 0.4777 - val_acc: 0.8816

Epoch 5/20

60000/60000 [=====] - 1s 24us/step - loss: 0.4852 -
acc: 0.8765 - val_loss: 0.4482 - val_acc: 0.8867

Epoch 6/20

60000/60000 [=====] - 1s 23us/step - loss: 0.4599 -
acc: 0.8813 - val_loss: 0.4267 - val_acc: 0.8901

Epoch 7/20

60000/60000 [=====] - 1s 24us/step - loss: 0.4410 -
acc: 0.8845 - val_loss: 0.4108 - val_acc: 0.8916

Epoch 8/20

60000/60000 [=====] - 1s 23us/step - loss: 0.4264 -
acc: 0.8873 - val_loss: 0.3983 - val_acc: 0.8955

Epoch 9/20

60000/60000 [=====] - 1s 24us/step - loss: 0.4145 -
acc: 0.8895 - val_loss: 0.3883 - val_acc: 0.8973

Epoch 10/20

60000/60000 [=====] - 1s 24us/step - loss: 0.4046 -
acc: 0.8911 - val_loss: 0.3792 - val_acc: 0.8997

Epoch 11/20

60000/60000 [=====] - 1s 23us/step - loss: 0.3964 -
acc: 0.8930 - val_loss: 0.3721 - val_acc: 0.9011

Epoch 12/20

60000/60000 [=====] - 1s 24us/step - loss: 0.3892 -
acc: 0.8950 - val_loss: 0.3659 - val_acc: 0.9026

Epoch 13/20

60000/60000 [=====] - 1s 24us/step - loss: 0.3829 -
acc: 0.8959 - val_loss: 0.3600 - val_acc: 0.9039

Epoch 14/20

60000/60000 [=====] - 1s 23us/step - loss: 0.3773 -
acc: 0.8973 - val_loss: 0.3551 - val_acc: 0.9053

Epoch 15/20

60000/60000 [=====] - 1s 23us/step - loss: 0.3724 -
acc: 0.8985 - val_loss: 0.3509 - val_acc: 0.9064

Epoch 16/20

60000/60000 [=====] - 1s 23us/step - loss: 0.3679 -
acc: 0.8993 - val_loss: 0.3471 - val_acc: 0.9072

Epoch 17/20

60000/60000 [=====] - 1s 23us/step - loss: 0.3639 -
acc: 0.9005 - val_loss: 0.3435 - val_acc: 0.9083

Epoch 18/20

60000/60000 [=====] - 1s 23us/step - loss: 0.3602 -
acc: 0.9014 - val_loss: 0.3400 - val_acc: 0.9087

Epoch 19/20

60000/60000 [=====] - 1s 23us/step - loss: 0.3568 -

acc: 0.9022 - val_loss: 0.3373 - val_acc: 0.9095

Epoch 20/20

60000/60000 [=====] - 1s 23us/step - loss: 0.3536 -

acc: 0.9029 - val_loss: 0.3345 - val_acc: 0.9097

```

In [0]: score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

#fig,ax = plt.subplots(1,1)
#ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))

# Test and train accuracy of the model
model_test_score = score[0]
model_test_acc = score[1]
model_train = history.history['acc']

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_loss : validation Loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to num
ber of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty)

```

Test score: 0.3344697316288948

Test accuracy: 0.9097



MLP + ReLu activation + Adam Optimizer + 2-Layer(experiment)

```
In [0]: # Multilayer perceptron

from keras.models import Sequential
from keras.layers import Dense, Activation
from keras.initializers import he_normal

#parameter
output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

model_sigmoid = Sequential()
# Adding first hidden layer
model_sigmoid.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))

# Adding second hidden layer
model_sigmoid.add(Dense(52, activation='relu', kernel_initializer=he_normal(seed=None)))

# Adding output layer
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()
```

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 364)	285740
dense_9 (Dense)	(None, 52)	18980
dense_10 (Dense)	(None, 10)	530
Total params: 305,250		
Trainable params: 305,250		
Non-trainable params: 0		


```
In [0]: model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history2 = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=n_b_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 38us/step - loss: 0.2754 -
acc: 0.9219 - val_loss: 0.1296 - val_acc: 0.9620

Epoch 2/20

60000/60000 [=====] - 2s 32us/step - loss: 0.1024 -
acc: 0.9696 - val_loss: 0.0875 - val_acc: 0.9737

Epoch 3/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0656 -
acc: 0.9799 - val_loss: 0.0865 - val_acc: 0.9747

Epoch 4/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0454 -
acc: 0.9863 - val_loss: 0.0758 - val_acc: 0.9754

Epoch 5/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0360 -
acc: 0.9891 - val_loss: 0.0660 - val_acc: 0.9788

Epoch 6/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0255 -
acc: 0.9923 - val_loss: 0.0702 - val_acc: 0.9785

Epoch 7/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0199 -
acc: 0.9940 - val_loss: 0.0767 - val_acc: 0.9790

Epoch 8/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0164 -
acc: 0.9948 - val_loss: 0.0718 - val_acc: 0.9808

Epoch 9/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0144 -
acc: 0.9955 - val_loss: 0.0730 - val_acc: 0.9799

Epoch 10/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0132 -
acc: 0.9956 - val_loss: 0.0789 - val_acc: 0.9801

Epoch 11/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0114 -
acc: 0.9963 - val_loss: 0.0866 - val_acc: 0.9781

Epoch 12/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0079 -
acc: 0.9978 - val_loss: 0.0837 - val_acc: 0.9787

Epoch 13/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0092 -
acc: 0.9969 - val_loss: 0.0771 - val_acc: 0.9803

Epoch 14/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0112 -
acc: 0.9961 - val_loss: 0.0821 - val_acc: 0.9787

Epoch 15/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0062 -
acc: 0.9981 - val_loss: 0.0826 - val_acc: 0.9806

Epoch 16/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0072 -
acc: 0.9977 - val_loss: 0.0812 - val_acc: 0.9808

Epoch 17/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0061 -
acc: 0.9980 - val_loss: 0.0807 - val_acc: 0.9813

Epoch 18/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0062 -
acc: 0.9979 - val_loss: 0.0970 - val_acc: 0.9785

Epoch 19/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0085 -

acc: 0.9969 - val_loss: 0.0989 - val_acc: 0.9790

Epoch 20/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0046 -

acc: 0.9986 - val_loss: 0.0922 - val_acc: 0.9814

```

In [0]: score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

#fig,ax = plt.subplots(1,1)
#ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
model_sigmoid_test_score = score[0]
model_sigmoid_test_acc = score[1]
model_sigmoid_train = history2.history['acc']

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to num
ber of epochs

vy = history2.history['val_loss']
ty = history2.history['loss']
plt_dynamic(x, vy, ty)

```

Test score: 0.09218703215732788

Test accuracy: 0.9814



```

In [0]: w_after = model_sigmoid.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + Sigmoid activation + ADAM

```
In [0]: model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 512)	401920
dense_6 (Dense)	(None, 128)	65664
dense_7 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 38us/step - loss: 0.5448 - acc: 0.8565 - val_loss: 0.2545 - val_acc: 0.9280

Epoch 2/20

60000/60000 [=====] - 2s 33us/step - loss: 0.2238 - acc: 0.9340 - val_loss: 0.1907 - val_acc: 0.9431

Epoch 3/20

60000/60000 [=====] - 2s 33us/step - loss: 0.1657 - acc: 0.9516 - val_loss: 0.1539 - val_acc: 0.9547

Epoch 4/20

60000/60000 [=====] - 2s 32us/step - loss: 0.1268 - acc: 0.9629 - val_loss: 0.1209 - val_acc: 0.9622

Epoch 5/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0989 - acc: 0.9709 - val_loss: 0.1023 - val_acc: 0.9682

Epoch 6/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0790 - acc: 0.9767 - val_loss: 0.0871 - val_acc: 0.9728

Epoch 7/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0648 - acc: 0.9810 - val_loss: 0.0787 - val_acc: 0.9749

Epoch 8/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0529 - acc: 0.9841 - val_loss: 0.0770 - val_acc: 0.9754

Epoch 9/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0424 - acc: 0.9871 - val_loss: 0.0717 - val_acc: 0.9777

Epoch 10/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0347 - acc: 0.9903 - val_loss: 0.0695 - val_acc: 0.9784

Epoch 11/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0279 - acc: 0.9923 - val_loss: 0.0629 - val_acc: 0.9801

Epoch 12/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0223 - acc: 0.9940 - val_loss: 0.0610 - val_acc: 0.9816

Epoch 13/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0187 - acc: 0.9950 - val_loss: 0.0669 - val_acc: 0.9806

Epoch 14/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0146 - acc: 0.9963 - val_loss: 0.0608 - val_acc: 0.9820

Epoch 15/20

```

60000/60000 [=====] - 2s 33us/step - loss: 0.0122 -
acc: 0.9973 - val_loss: 0.0683 - val_acc: 0.9796
Epoch 16/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0097 -
acc: 0.9977 - val_loss: 0.0676 - val_acc: 0.9792
Epoch 17/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0074 -
acc: 0.9984 - val_loss: 0.0696 - val_acc: 0.9796
Epoch 18/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0065 -
acc: 0.9987 - val_loss: 0.0734 - val_acc: 0.9778
Epoch 19/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0055 -
acc: 0.9988 - val_loss: 0.0775 - val_acc: 0.9794
Epoch 20/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0052 -
acc: 0.9986 - val_loss: 0.0684 - val_acc: 0.9815

```

```

In [0]: score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to num
ber of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.06838007308253509
Test accuracy: 0.9815

```



```
In [0]: w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

MLP + ReLU +SGD

```

In [0]: # Multilayer perceptron

# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0, \sigma) = N(0, 0.062)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0, \sigma) = N(0, 0.125)$ 
# out =>  $\sigma = \sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()

```

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 512)	401920
dense_9 (Dense)	(None, 128)	65664
dense_10 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

```
In [0]: model_relu.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=[  
        'accuracy'])  
  
        history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_ep  
        och, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 32us/step - loss: 0.7739 -
acc: 0.7814 - val_loss: 0.3893 - val_acc: 0.8934

Epoch 2/20

60000/60000 [=====] - 2s 28us/step - loss: 0.3523 -
acc: 0.9004 - val_loss: 0.2989 - val_acc: 0.9156

Epoch 3/20

60000/60000 [=====] - 2s 28us/step - loss: 0.2888 -
acc: 0.9170 - val_loss: 0.2611 - val_acc: 0.9269

Epoch 4/20

60000/60000 [=====] - 2s 28us/step - loss: 0.2540 -
acc: 0.9280 - val_loss: 0.2368 - val_acc: 0.9321

Epoch 5/20

60000/60000 [=====] - 2s 28us/step - loss: 0.2299 -
acc: 0.9343 - val_loss: 0.2197 - val_acc: 0.9369

Epoch 6/20

60000/60000 [=====] - 2s 28us/step - loss: 0.2116 -
acc: 0.9396 - val_loss: 0.2043 - val_acc: 0.9401

Epoch 7/20

60000/60000 [=====] - 2s 28us/step - loss: 0.1965 -
acc: 0.9439 - val_loss: 0.1924 - val_acc: 0.9433

Epoch 8/20

60000/60000 [=====] - 2s 28us/step - loss: 0.1841 -
acc: 0.9470 - val_loss: 0.1826 - val_acc: 0.9459

Epoch 9/20

60000/60000 [=====] - 2s 28us/step - loss: 0.1734 -
acc: 0.9506 - val_loss: 0.1744 - val_acc: 0.9477

Epoch 10/20

60000/60000 [=====] - 2s 28us/step - loss: 0.1641 -
acc: 0.9531 - val_loss: 0.1678 - val_acc: 0.9489

Epoch 11/20

60000/60000 [=====] - 2s 28us/step - loss: 0.1559 -
acc: 0.9555 - val_loss: 0.1607 - val_acc: 0.9524

Epoch 12/20

60000/60000 [=====] - 2s 28us/step - loss: 0.1483 -
acc: 0.9581 - val_loss: 0.1550 - val_acc: 0.9528

Epoch 13/20

60000/60000 [=====] - 2s 29us/step - loss: 0.1420 -
acc: 0.9593 - val_loss: 0.1496 - val_acc: 0.9552

Epoch 14/20

60000/60000 [=====] - 2s 28us/step - loss: 0.1358 -
acc: 0.9612 - val_loss: 0.1445 - val_acc: 0.9576

Epoch 15/20

60000/60000 [=====] - 2s 28us/step - loss: 0.1302 -
acc: 0.9624 - val_loss: 0.1397 - val_acc: 0.9589

Epoch 16/20

60000/60000 [=====] - 2s 28us/step - loss: 0.1250 -
acc: 0.9643 - val_loss: 0.1372 - val_acc: 0.9594

Epoch 17/20

60000/60000 [=====] - 2s 28us/step - loss: 0.1203 -
acc: 0.9658 - val_loss: 0.1326 - val_acc: 0.9605

Epoch 18/20

60000/60000 [=====] - 2s 28us/step - loss: 0.1160 -
acc: 0.9669 - val_loss: 0.1304 - val_acc: 0.9618

Epoch 19/20

60000/60000 [=====] - 2s 28us/step - loss: 0.1119 -

```
acc: 0.9686 - val_loss: 0.1260 - val_acc: 0.9623
Epoch 20/20
60000/60000 [=====] - 2s 28us/step - loss: 0.1082 -
acc: 0.9693 - val_loss: 0.1236 - val_acc: 0.9636
```

```
In [0]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to num
ber of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

```
Test score: 0.1236097104612738
Test accuracy: 0.9636
```

```
In [0]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

MLP + ReLU + ADAM

```
In [0]: model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_
initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(m
ean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=
['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_ep
och, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 512)	401920
dense_12 (Dense)	(None, 128)	65664
dense_13 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 39us/step - loss: 0.2362 - acc: 0.9292 - val_loss: 0.1135 - val_acc: 0.9645

Epoch 2/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0872 - acc: 0.9743 - val_loss: 0.0889 - val_acc: 0.9730

Epoch 3/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0525 - acc: 0.9837 - val_loss: 0.0720 - val_acc: 0.9768

Epoch 4/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0363 - acc: 0.9892 - val_loss: 0.0656 - val_acc: 0.9783

Epoch 5/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0276 - acc: 0.9917 - val_loss: 0.0638 - val_acc: 0.9805

Epoch 6/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0184 - acc: 0.9942 - val_loss: 0.0703 - val_acc: 0.9802

Epoch 7/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0172 - acc: 0.9948 - val_loss: 0.0765 - val_acc: 0.9782

Epoch 8/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0155 - acc: 0.9953 - val_loss: 0.0747 - val_acc: 0.9785

Epoch 9/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0118 - acc: 0.9962 - val_loss: 0.0852 - val_acc: 0.9765

Epoch 10/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0131 - acc: 0.9958 - val_loss: 0.0767 - val_acc: 0.9811

Epoch 11/20

60000/60000 [=====] - 2s 34us/step - loss: 0.0107 - acc: 0.9964 - val_loss: 0.0913 - val_acc: 0.9771

Epoch 12/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0096 - acc: 0.9967 - val_loss: 0.1016 - val_acc: 0.9763

Epoch 13/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0125 - acc: 0.9958 - val_loss: 0.0791 - val_acc: 0.9797

Epoch 14/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0055 - acc: 0.9982 - val_loss: 0.0825 - val_acc: 0.9813


```

Epoch 15/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0112 -
acc: 0.9962 - val_loss: 0.0926 - val_acc: 0.9783
Epoch 16/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0092 -
acc: 0.9970 - val_loss: 0.0891 - val_acc: 0.9808
Epoch 17/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0072 -
acc: 0.9975 - val_loss: 0.0898 - val_acc: 0.9806
Epoch 18/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0030 -
acc: 0.9991 - val_loss: 0.0939 - val_acc: 0.9818
Epoch 19/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0060 -
acc: 0.9982 - val_loss: 0.1058 - val_acc: 0.9789
Epoch 20/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0134 -
acc: 0.9956 - val_loss: 0.0915 - val_acc: 0.9811

```

```

In [0]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to num
ber of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09147267398560358

Test accuracy: 0.9811

```
In [0]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

MLP + Batch-Norm on hidden Layers + AdamOptimizer </2>

```

In [0]: #type of initializer ,,https://keras.io/initializers/
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=\sqrt{2/(n_i+n_{i+1})}$ .
# h1 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.039 \Rightarrow N(0,\sigma) = N(0,0.039)$ 
# h2 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.055 \Rightarrow N(0,\sigma) = N(0,0.055)$ 
# h1 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```

Layer (type)	Output Shape	Param #
dense_14 (Dense)	(None, 512)	401920
batch_normalization_1 (Batch Normalization)	(None, 512)	2048
dense_15 (Dense)	(None, 128)	65664
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dense_16 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

```
In [0]: model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics
      =['accuracy'])

      history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_e
      poch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 68us/step - loss: 0.3067 -
acc: 0.9101 - val_loss: 0.2085 - val_acc: 0.9374

Epoch 2/20

60000/60000 [=====] - 3s 55us/step - loss: 0.1726 -
acc: 0.9494 - val_loss: 0.1664 - val_acc: 0.9507

Epoch 3/20

60000/60000 [=====] - 3s 56us/step - loss: 0.1375 -
acc: 0.9593 - val_loss: 0.1487 - val_acc: 0.9547

Epoch 4/20

60000/60000 [=====] - 3s 55us/step - loss: 0.1116 -
acc: 0.9666 - val_loss: 0.1414 - val_acc: 0.9571

Epoch 5/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0939 -
acc: 0.9712 - val_loss: 0.1264 - val_acc: 0.9610

Epoch 6/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0828 -
acc: 0.9751 - val_loss: 0.1172 - val_acc: 0.9662

Epoch 7/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0703 -
acc: 0.9780 - val_loss: 0.1144 - val_acc: 0.9666

Epoch 8/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0596 -
acc: 0.9816 - val_loss: 0.1085 - val_acc: 0.9650

Epoch 9/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0521 -
acc: 0.9834 - val_loss: 0.1033 - val_acc: 0.9671

Epoch 10/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0458 -
acc: 0.9860 - val_loss: 0.1072 - val_acc: 0.9690

Epoch 11/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0395 -
acc: 0.9876 - val_loss: 0.1001 - val_acc: 0.9703

Epoch 12/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0375 -
acc: 0.9879 - val_loss: 0.1070 - val_acc: 0.9696

Epoch 13/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0321 -
acc: 0.9898 - val_loss: 0.1056 - val_acc: 0.9702

Epoch 14/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0272 -
acc: 0.9911 - val_loss: 0.0940 - val_acc: 0.9733

Epoch 15/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0247 -
acc: 0.9918 - val_loss: 0.1031 - val_acc: 0.9723

Epoch 16/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0223 -
acc: 0.9926 - val_loss: 0.0939 - val_acc: 0.9737

Epoch 17/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0200 -
acc: 0.9932 - val_loss: 0.1072 - val_acc: 0.9721

Epoch 18/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0205 -
acc: 0.9929 - val_loss: 0.1022 - val_acc: 0.9736

Epoch 19/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0178 -

acc: 0.9941 - val_loss: 0.0966 - val_acc: 0.9744
 Epoch 20/20
 60000/60000 [=====] - 3s 55us/step - loss: 0.0154 -
 acc: 0.9949 - val_loss: 0.0993 - val_acc: 0.9745

```
In [0]: score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to num
ber of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09926731985980877
 Test accuracy: 0.9745

```

In [0]: w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

```

In [0]: # %matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, y, y_1, ax, ticks,title, colors=['b']):
    ax.plot(x, y, 'b', label="Train Loss")
    ax.plot(x, y_1, 'r', label="Test Loss")
    if len(x)==1:
        plt.legend()
        plt.title(title)
    plt.yticks(ticks)
    fig.canvas.draw()

```

MLP + ReLu activation + Adam Optimizer + BN+Dropout+ 2-Layer

```
In [0]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization

from keras.models import Sequential
from keras.layers import Dense, Activation
from keras.initializers import he_normal

#parameter
output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

# Initialising model
model_drop = Sequential()

# Adding first hidden layer
model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model_drop.add(BatchNormalization())
# Adding dropout to first hidden layer
model_drop.add(Dropout(0.5))

# Adding second hidden layer
model_drop.add(Dense(52, activation='relu', kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model_drop.add(BatchNormalization())
# Adding dropout to second hidden layer
model_drop.add(Dropout(0.5))

# Adding output layer
model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```


Layer (type)	Output Shape	Param #
=====	=====	=====
dense_27 (Dense)	(None, 364)	285740
batch_normalization_13 (Batch Normalization)	(None, 364)	1456
dropout_13 (Dropout)	(None, 364)	0
dense_28 (Dense)	(None, 52)	18980
batch_normalization_14 (Batch Normalization)	(None, 52)	208
dropout_14 (Dropout)	(None, 52)	0
dense_29 (Dense)	(None, 10)	530
=====	=====	=====
Total params: 306,914		
Trainable params: 306,082		
Non-trainable params: 832		
=====		

```
In [0]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=
        ['accuracy'])

        history3 = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_e
        poch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 82us/step - loss: 0.5326 -
acc: 0.8427 - val_loss: 0.1750 - val_acc: 0.9450

Epoch 2/20

60000/60000 [=====] - 3s 57us/step - loss: 0.2624 -
acc: 0.9245 - val_loss: 0.1216 - val_acc: 0.9632

Epoch 3/20

60000/60000 [=====] - 3s 57us/step - loss: 0.2014 -
acc: 0.9417 - val_loss: 0.0987 - val_acc: 0.9686

Epoch 4/20

60000/60000 [=====] - 3s 57us/step - loss: 0.1738 -
acc: 0.9502 - val_loss: 0.0906 - val_acc: 0.9715

Epoch 5/20

60000/60000 [=====] - 3s 57us/step - loss: 0.1544 -
acc: 0.9547 - val_loss: 0.0848 - val_acc: 0.9737

Epoch 6/20

60000/60000 [=====] - 3s 57us/step - loss: 0.1387 -
acc: 0.9591 - val_loss: 0.0833 - val_acc: 0.9743

Epoch 7/20

60000/60000 [=====] - 3s 57us/step - loss: 0.1307 -
acc: 0.9629 - val_loss: 0.0802 - val_acc: 0.9760

Epoch 8/20

60000/60000 [=====] - 3s 57us/step - loss: 0.1232 -
acc: 0.9643 - val_loss: 0.0814 - val_acc: 0.9760

Epoch 9/20

60000/60000 [=====] - 3s 57us/step - loss: 0.1133 -
acc: 0.9672 - val_loss: 0.0784 - val_acc: 0.9774

Epoch 10/20

60000/60000 [=====] - 3s 58us/step - loss: 0.1072 -
acc: 0.9682 - val_loss: 0.0677 - val_acc: 0.9804

Epoch 11/20

60000/60000 [=====] - 3s 57us/step - loss: 0.1039 -
acc: 0.9696 - val_loss: 0.0672 - val_acc: 0.9783

Epoch 12/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0975 -
acc: 0.9709 - val_loss: 0.0728 - val_acc: 0.9770

Epoch 13/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0944 -
acc: 0.9718 - val_loss: 0.0699 - val_acc: 0.9789

Epoch 14/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0887 -
acc: 0.9732 - val_loss: 0.0710 - val_acc: 0.9795

Epoch 15/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0876 -
acc: 0.9734 - val_loss: 0.0666 - val_acc: 0.9809

Epoch 16/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0811 -
acc: 0.9758 - val_loss: 0.0717 - val_acc: 0.9797

Epoch 17/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0797 -
acc: 0.9761 - val_loss: 0.0697 - val_acc: 0.9793

Epoch 18/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0802 -
acc: 0.9765 - val_loss: 0.0631 - val_acc: 0.9811

Epoch 19/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0765 -

acc: 0.9773 - val_loss: 0.0617 - val_acc: 0.9813

Epoch 20/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0737 -

acc: 0.9780 - val_loss: 0.0639 - val_acc: 0.9810

```
In [0]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

#fig,ax = plt.subplots(1,1)
#ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# Test and train accuracy of the model
model_drop_test_score = score[0]
model_drop_test_acc = score[1]
model_drop_train = history3.history['acc']

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
epoch, verbose=1, validation_data=(X_test, Y_test))

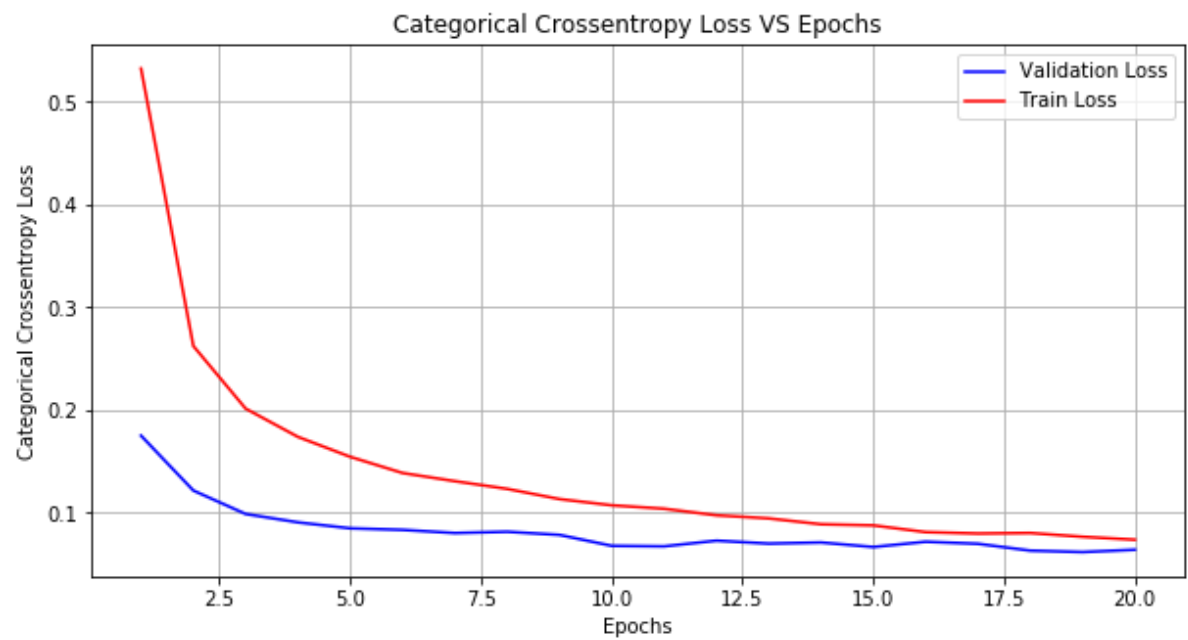
# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to num
ber of epochs

vy = history3.history['val_loss']
ty = history3.history['loss']
plt_dynamic(x, vy, ty)
```

Test score: 0.06392139566795085

Test accuracy: 0.981



```

In [0]: w_after = model_drop.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + ReLu activation + Adam Optimizer + BN+Dropout+ 3-Layer

```
In [0]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization

from keras.models import Sequential
from keras.layers import Dense, Activation
from keras.initializers import he_normal

#parameter
output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

# Initialising model
model3_drop = Sequential()

# Adding first hidden layer
model3_drop.add(Dense(554, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
model3_drop.add(BatchNormalization())
model3_drop.add(Dropout(0.5))

# Adding second hidden layer
model3_drop.add(Dense(225, activation='relu', kernel_initializer=he_normal(seed=None)))
model3_drop.add(BatchNormalization())
model3_drop.add(Dropout(0.5))

# Adding third hidden layer
model3_drop.add(Dense(78, activation='relu', kernel_initializer=he_normal(seed=None)))
model3_drop.add(BatchNormalization())
model3_drop.add(Dropout(0.5))

model3_drop.add(Dense(output_dim, activation='softmax'))

model3_drop.summary()
```


Layer (type)	Output Shape	Param #
dense_30 (Dense)	(None, 554)	434890
dense_31 (Dense)	(None, 225)	124875
dense_32 (Dense)	(None, 78)	17628
batch_normalization_17 (Batch Normalization)	(None, 78)	312
dropout_17 (Dropout)	(None, 78)	0
dense_33 (Dense)	(None, 10)	790
Total params: 578,495		
Trainable params: 578,339		
Non-trainable params: 156		

```
In [0]: model3_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics
      =['accuracy'])

      history4 = model3_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
      epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 73us/step - loss: 0.3175 -
acc: 0.9073 - val_loss: 0.1184 - val_acc: 0.9651

Epoch 2/20

60000/60000 [=====] - 3s 50us/step - loss: 0.1226 -
acc: 0.9649 - val_loss: 0.0953 - val_acc: 0.9699

Epoch 3/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0795 -
acc: 0.9769 - val_loss: 0.0917 - val_acc: 0.9740

Epoch 4/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0608 -
acc: 0.9818 - val_loss: 0.0717 - val_acc: 0.9785

Epoch 5/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0457 -
acc: 0.9858 - val_loss: 0.0913 - val_acc: 0.9734

Epoch 6/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0399 -
acc: 0.9879 - val_loss: 0.0932 - val_acc: 0.9738

Epoch 7/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0302 -
acc: 0.9910 - val_loss: 0.0754 - val_acc: 0.9791

Epoch 8/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0252 -
acc: 0.9921 - val_loss: 0.1031 - val_acc: 0.9734

Epoch 9/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0237 -
acc: 0.9926 - val_loss: 0.0623 - val_acc: 0.9834

Epoch 10/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0190 -
acc: 0.9943 - val_loss: 0.0746 - val_acc: 0.9799

Epoch 11/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0188 -
acc: 0.9943 - val_loss: 0.0799 - val_acc: 0.9803

Epoch 12/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0169 -
acc: 0.9946 - val_loss: 0.0854 - val_acc: 0.9801

Epoch 13/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0178 -
acc: 0.9942 - val_loss: 0.0702 - val_acc: 0.9829

Epoch 14/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0139 -
acc: 0.9957 - val_loss: 0.0830 - val_acc: 0.9817

Epoch 15/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0138 -
acc: 0.9956 - val_loss: 0.0826 - val_acc: 0.9810

Epoch 16/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0106 -
acc: 0.9966 - val_loss: 0.0690 - val_acc: 0.9850

Epoch 17/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0091 -
acc: 0.9971 - val_loss: 0.0845 - val_acc: 0.9812

Epoch 18/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0140 -
acc: 0.9955 - val_loss: 0.0850 - val_acc: 0.9810

Epoch 19/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0090 -

acc: 0.9972 - val_loss: 0.0732 - val_acc: 0.9849

Epoch 20/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0092 -

acc: 0.9972 - val_loss: 0.0766 - val_acc: 0.9838

```

In [0]: score = model3_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

#fig,ax = plt.subplots(1,1)
#ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# Test and train accuracy of the model
model3_drop_test_score = score[0]
model3_drop_test_acc = score[1]
model3_drop_train = history4.history['acc']

# List of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to num
ber of epochs

vy = history4.history['val_loss']
ty = history4.history['loss']
plt_dynamic(x, vy, ty)

```

Test score: 0.07663734175693944

Test accuracy: 0.9838



MLP + ReLu activation + Adam Optimizer + BN+Dropout+ 5-Layer

```
In [0]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization

from keras.models import Sequential
from keras.layers import Dense, Activation
from keras.initializers import he_normal

#paramter
output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

# Initialising model
model4_drop = Sequential()

# Adding first hidden layer
model4_drop.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model4_drop.add(BatchNormalization())
# Adding dropout
model4_drop.add(Dropout(0.5))

# Adding second hidden layer
model4_drop.add(Dense(256, activation='relu', kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model4_drop.add(BatchNormalization())
# Adding dropout
model4_drop.add(Dropout(0.5))

# Adding third hidden layer
model4_drop.add(Dense(128, activation='relu', kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model4_drop.add(BatchNormalization())
# Adding dropout
model4_drop.add(Dropout(0.5))

# Adding fourth hidden layer
model4_drop.add(Dense(64, activation='relu', kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model4_drop.add(BatchNormalization())
# Adding dropout
model4_drop.add(Dropout(0.5))

# Adding fifth hidden layer
model4_drop.add(Dense(32, activation='relu', kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
```

```

model4_drop.add(BatchNormalization())
# Adding dropout
model4_drop.add(Dropout(0.5))

# Adding output layer
model4_drop.add(Dense(output_dim, activation='softmax'))

model4_drop.summary()

```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_34 (Dense)	(None, 512)	401920
batch_normalization_18 (Batch Normalization)	(None, 512)	2048
dropout_18 (Dropout)	(None, 512)	0
dense_35 (Dense)	(None, 256)	131328
batch_normalization_19 (Batch Normalization)	(None, 256)	1024
dropout_19 (Dropout)	(None, 256)	0
dense_36 (Dense)	(None, 128)	32896
batch_normalization_20 (Batch Normalization)	(None, 128)	512
dropout_20 (Dropout)	(None, 128)	0
dense_37 (Dense)	(None, 64)	8256
batch_normalization_21 (Batch Normalization)	(None, 64)	256
dropout_21 (Dropout)	(None, 64)	0
dense_38 (Dense)	(None, 32)	2080
batch_normalization_22 (Batch Normalization)	(None, 32)	128
dropout_22 (Dropout)	(None, 32)	0
dense_39 (Dense)	(None, 10)	330
=====	=====	=====
Total params: 580,778		
Trainable params: 578,794		
Non-trainable params: 1,984		
=====	=====	=====


```
In [0]: model4_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics
      =['accuracy'])

      history5 = model4_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
      epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 138us/step - loss: 1.4364 -
acc: 0.5277 - val_loss: 0.3345 - val_acc: 0.9116

Epoch 2/20

60000/60000 [=====] - 6s 99us/step - loss: 0.5814 -
acc: 0.8349 - val_loss: 0.1847 - val_acc: 0.9483

Epoch 3/20

60000/60000 [=====] - 6s 99us/step - loss: 0.3939 -
acc: 0.8993 - val_loss: 0.1560 - val_acc: 0.9589

Epoch 4/20

60000/60000 [=====] - 6s 98us/step - loss: 0.3255 -
acc: 0.9214 - val_loss: 0.1461 - val_acc: 0.9617

Epoch 5/20

60000/60000 [=====] - 6s 98us/step - loss: 0.2819 -
acc: 0.9328 - val_loss: 0.1263 - val_acc: 0.9671

Epoch 6/20

60000/60000 [=====] - 6s 99us/step - loss: 0.2521 -
acc: 0.9405 - val_loss: 0.1182 - val_acc: 0.9713

Epoch 7/20

60000/60000 [=====] - 6s 100us/step - loss: 0.2322 -
acc: 0.9458 - val_loss: 0.1061 - val_acc: 0.9732

Epoch 8/20

60000/60000 [=====] - 6s 103us/step - loss: 0.2182 -
acc: 0.9490 - val_loss: 0.1054 - val_acc: 0.9738

Epoch 9/20

60000/60000 [=====] - 6s 101us/step - loss: 0.2002 -
acc: 0.9541 - val_loss: 0.1061 - val_acc: 0.9751

Epoch 10/20

60000/60000 [=====] - 6s 100us/step - loss: 0.1920 -
acc: 0.9565 - val_loss: 0.0983 - val_acc: 0.9764

Epoch 11/20

60000/60000 [=====] - 6s 100us/step - loss: 0.1820 -
acc: 0.9591 - val_loss: 0.0862 - val_acc: 0.9783

Epoch 12/20

60000/60000 [=====] - 6s 99us/step - loss: 0.1715 -
acc: 0.9614 - val_loss: 0.0839 - val_acc: 0.9794

Epoch 13/20

60000/60000 [=====] - 6s 99us/step - loss: 0.1711 -
acc: 0.9610 - val_loss: 0.0848 - val_acc: 0.9797

Epoch 14/20

60000/60000 [=====] - 6s 99us/step - loss: 0.1567 -
acc: 0.9649 - val_loss: 0.0895 - val_acc: 0.9785

Epoch 15/20

60000/60000 [=====] - 6s 99us/step - loss: 0.1546 -
acc: 0.9650 - val_loss: 0.0815 - val_acc: 0.9814

Epoch 16/20

60000/60000 [=====] - 6s 98us/step - loss: 0.1456 -
acc: 0.9668 - val_loss: 0.0877 - val_acc: 0.9802

Epoch 17/20

60000/60000 [=====] - 6s 99us/step - loss: 0.1425 -
acc: 0.9676 - val_loss: 0.0804 - val_acc: 0.9811

Epoch 18/20

60000/60000 [=====] - 6s 99us/step - loss: 0.1403 -
acc: 0.9682 - val_loss: 0.0761 - val_acc: 0.9825

Epoch 19/20

60000/60000 [=====] - 6s 99us/step - loss: 0.1369 -

acc: 0.9696 - val_loss: 0.0788 - val_acc: 0.9816

Epoch 20/20

60000/60000 [=====] - 6s 101us/step - loss: 0.1322 -

acc: 0.9698 - val_loss: 0.0812 - val_acc: 0.9820

```

In [0]: score = model4_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

#fig,ax = plt.subplots(1,1)
#ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
model4_drop_test_score = score[0]
model4_drop_test_acc = score[1]
model4_drop_train = history5.history['acc']

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to num
ber of epochs

vy = history5.history['val_loss']
ty = history5.history['loss']
plt_dynamic(x, vy, ty)

```

Test score: 0.08122210750945379

Test accuracy: 0.982



Summarize

```
In [0]: from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["Model", "Train accuracy", "Test accuracy"]
x.add_row(["MLP + ReLu activation + Adam Optimizer + BN+Dropout+ 2-Layer", model_drop_train[19], model_drop_test_acc])
x.add_row(["MLP + ReLu activation + Adam Optimizer + BN+Dropout+ 3-Layer", model3_drop_train[19], model3_drop_test_acc])
x.add_row(["MLP + ReLu activation + Adam Optimizer + BN+Dropout+ 5-Layer", model4_drop_train[19], model4_drop_test_acc])
print(x)
```

Test accuracy		Model	Train accuracy
651224	0.981	MLP + ReLu activation + Adam Optimizer + BN+Dropout+ 2-Layer	0.9780333333
333333	0.9838	MLP + ReLu activation + Adam Optimizer + BN+Dropout+ 3-Layer	0.9971833333
348775	0.982	MLP + ReLu activation + Adam Optimizer + BN+Dropout+ 5-Layer	0.9697666666

Observation:

for 2-Layer Model, train accuracy little lower than test accuracy

train accuracy also low for higher number of layer

Test accuracy almost same for all case